

Building a machine learning model to predict product back-orders

EDSA Bootcamp 2018

André Águas
António Correia
João Januário
Valter Bento

M20170973@novaims.unl.pt
M20170975@novaims.unl.pt
M20170985@novaims.unl.pt
M20170999@novaims.unl.pt

AGENDA




Introduction



Exploratory data analysis



Data pre-processing



Feature analysis



Model pipeline and evaluation metrics



Model results and proof of concept



Conclusions

INTRODUCTION



What?

- BI4all
- Inventory Management
- Predicting Back Orders

Why?

- "Stock-Out"
- Signals fast growing companies
- Reputational risk
- Risk to lose clients
- Inventory Risk (spoilage, theft, damage)

INTRODUCTION



How can we increase profit?

Excess inventory
(increased
storage costs)

Accurate
predictions of stock
needs and back-
orders occurrences

Unable to predict
"stock-out"
(increased cost of
lost sales)

EXPLORATORY DATA ANALYSIS 🔍

- 8-weeks historical data
- Predict next week

Data types

- Current inventory
- Forecast
- Sales
- Performance

```
In [11]: df.shape
```

```
Out[11]: (1687861, 23)
```

```
In [9]: df.dtypes
```

```
Out[9]: sku                object
national_inv             float64
lead_time                float64
in_transit_qty           float64
forecast_3_month         float64
forecast_6_month         float64
forecast_9_month         float64
sales_1_month            float64
sales_3_month            float64
sales_6_month            float64
sales_9_month            float64
min_bank                 float64
potential_issue           object
pieces_past_due          float64
perf_6_month_avg         float64
perf_12_month_avg        float64
local_bo_qty             float64
deck_risk                object
oe_constraint            object
ppap_risk                object
stop_auto_buy            object
rev_stop                 object
went_on_backorder        object
dtype: object
```

EXPLORATORY DATA ANALYSIS

```
In [10]: df.isnull().sum()/len(df)*100
```

```
Out[10]: sku                0.000000
national_inv              0.000059
lead_time                5.977625
in_transit_qty           0.000059
forecast_3_month         0.000059
forecast_6_month         0.000059
forecast_9_month         0.000059
sales_1_month            0.000059
sales_3_month            0.000059
sales_6_month            0.000059
sales_9_month            0.000059
min_bank                 0.000059
potential_issue          0.000059
pieces_past_due          0.000059
perf_6_month_avg         0.000059
perf_12_month_avg        0.000059
local_bo_qty             0.000059
deck_risk                0.000059
oe_constraint            0.000059
ppap_risk                0.000059
stop_auto_buy            0.000059
rev_stop                 0.000059
went_on_backorder        0.000059
dtype: float64
```

```
In [14]: df.describe().T
```

```
Out[14]:
```


	count	mean	std	min	25%	50%	75%	max
national_inv	1687860.0	496.111782	29615.233831	-27256.0	4.00	15.00	80.00	12334404.0
lead_time	1586967.0	7.872267	7.056024	0.0	4.00	8.00	9.00	52.0
in_transit_qty	1687860.0	44.052022	1342.741731	0.0	0.00	0.00	0.00	489408.0
forecast_3_month	1687860.0	178.119284	5026.553102	0.0	0.00	0.00	4.00	1427612.0
forecast_6_month	1687860.0	344.986664	9795.151861	0.0	0.00	0.00	12.00	2461360.0
forecast_9_month	1687860.0	506.364431	14378.923562	0.0	0.00	0.00	20.00	3777304.0
sales_1_month	1687860.0	55.926069	1928.195879	0.0	0.00	0.00	4.00	741774.0
sales_3_month	1687860.0	175.025930	5192.377625	0.0	0.00	1.00	15.00	1105478.0
sales_6_month	1687860.0	341.728839	9613.167104	0.0	0.00	2.00	31.00	2146625.0
sales_9_month	1687860.0	525.269701	14838.613523	0.0	0.00	4.00	47.00	3205172.0
min_bank	1687860.0	52.772303	1254.983089	0.0	0.00	0.00	3.00	313319.0
pieces_past_due	1687860.0	2.043724	236.016500	0.0	0.00	0.00	0.00	146496.0
perf_6_month_avg	1687860.0	-6.872059	26.556357	-99.0	0.63	0.82	0.97	1.0
perf_12_month_avg	1687860.0	-6.437947	25.843331	-99.0	0.66	0.81	0.95	1.0
local_bo_qty	1687860.0	0.626451	33.722242	0.0	0.00	0.00	0.00	12530.0

DATA PRE-PROCESSING



- Data Conversion [String → Numeric] (“Yes/No” → “1/0”)
- Drop lines with NULL values (↓ 6% train data)
- Replace performance “NULL values” (-99) with median:
 - perf_6_month_avg (≈ 2% train data)
 - perf_12_month_avg (≈ 1% train data)
- Normalization


FEATURE ANALYSIS



Target variable - went_on_backorder

Despite having more than 1.5M observations on our set, we only have 10k observations classified as 1, less than 1%.

The small number of observations classified as 1's might be a challenge during training process of the models




went_on_backorder		
0	1,575,998	99.3%
1	10,969	0.7%


Relation between independent variables

Applying a correlation to all independent variables we determine a strong relation between three groups

- Forecast_3_month, forecast_6_month and forecast_9_month
- Sales_1_month, Sales_3_month, sales_6_month and sales_9_month
- Perf_6_month_avg and Perf_12_month_avg



FEATURE ANALYSIS




Relation between independent variables

Beside correlation between variables we performed an Random Forest model to assess the importance of a variable to the model accuracy and we compared the Mean and Std Deviation when went_on_backorder is 0 and 1

variables	RF Model importance	Mean		Std Dev	
		0	1	0	1
national_inv	3.54	499	21	29,715	608
forecast_3_month	1.43	178	157	5,042	1,635
forecast_6_month	1.08	346	245	9,826	2,457
forecast_9_month	0.58	508	326	14,425	3,145
in_transit_qty	0.50	44	4	1,347	47
sales_1_month	0.50	56	29	1,935	273
sales_3_month	0.41	176	79	5,210	509
stop_auto_buy	0.30	1	1	0	0
local_bo_qty	0.28	1	5	33	62
sales_6_month	0.24	343	139	9,645	901
sales_9_month	0.21	527	206	14,888	1,375
lead_time	0.19	6,006	2,875	23,744	16,693
perf_6_month_avg	0.19	-7	-3	27	19
perf_12_month_avg	0.18	-6	-3	26	18
min_bank	0.15	53	24	1,259	145
pieces_past_due	0.09	2	4	237	34
ppap_risk	0.08	0	0	0	0
deck_risk	0.07	0	0	0	0
potential_issue	0.00	0	0	0	0
oe_constraint	0.00	0	0	0	0
rev_stop	0.00	0	0	0	0

FEATURE ANALYSIS



Selecting best variables for our model

We have selected 11 of 21 variables as relevant features for use in model construction.

We have taken into account the correlation between variables, the importance according to the Random Forest and Mean and Std Deviation analysis

Initial Variables

national_inv
forecast_3_month
forecast_6_month
forecast_9_month
in_transit_qty
sales_1_month
sales_3_month
stop_auto_buy
local_bo_qty
sales_6_month
sales_9_month
lead_time
perf_6_month_avg
perf_12_month_avg
min_bank
pieces_past_due
ppap_risk
deck_risk
potential_issue
oe_constraint
rev_stop

Final Variables

national_inv
forecast_3_month

in_transit_qty
sales_1_month


stop_auto_buy
local_bo_qty

lead_time
perf_6_month_avg

min_bank
pieces_past_due

potential_issue


MODEL PIPELINE AND EVALUATION



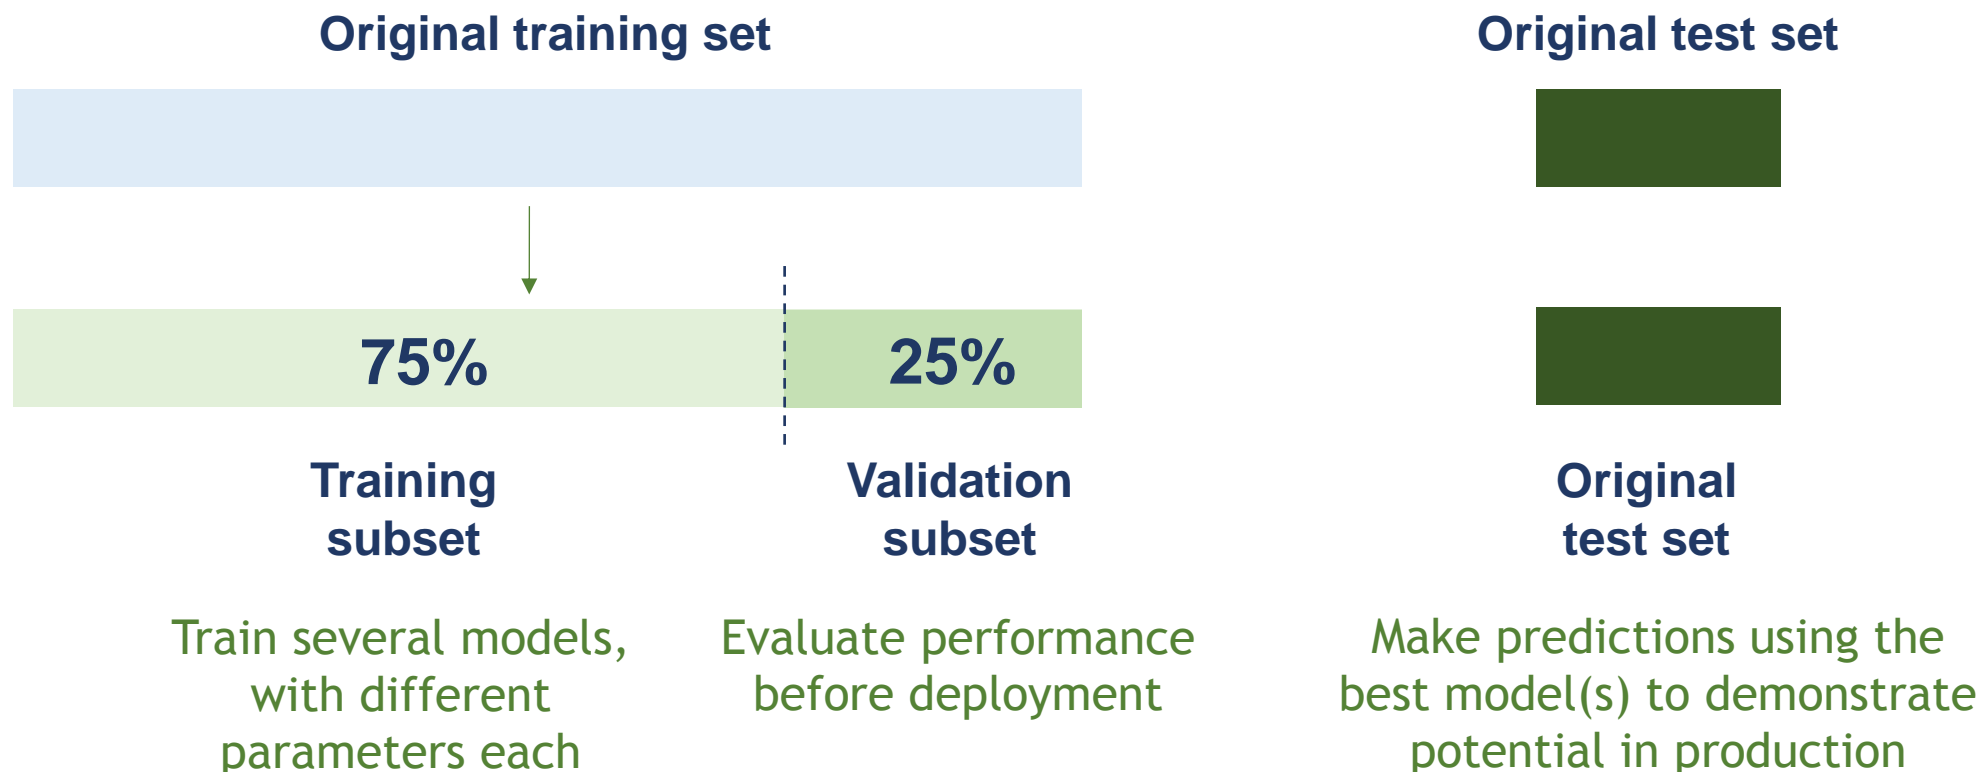
Selecting and adjusting the best models: avoiding model overfitting

- The ‘Test’ dataset will be used to create a proof of concept, hence we do not want to adjust models based on it in order to avoid overfitting
- We will thus assess models and tweak their parameters by randomly splitting the ‘Train’ dataset into a Train and Validation subsets
- Then we can use the new Train subset to train different models with varying configurations and use the Validation set to check their performance


MODEL PIPELINE AND EVALUATION



Selecting and adjusting the best models: avoiding model overfitting




MODEL PIPELINE AND EVALUATION




Addressing data imbalance

- Our dataset is extremely imbalanced, and we address this issue with an over-sampling algorithm
- SMOTE artificially creates observations based on the K nearest neighbors of the minority class observations
- To use SMOTE correctly, we should apply it only to our training subset (i.e., after splitting the original training set)

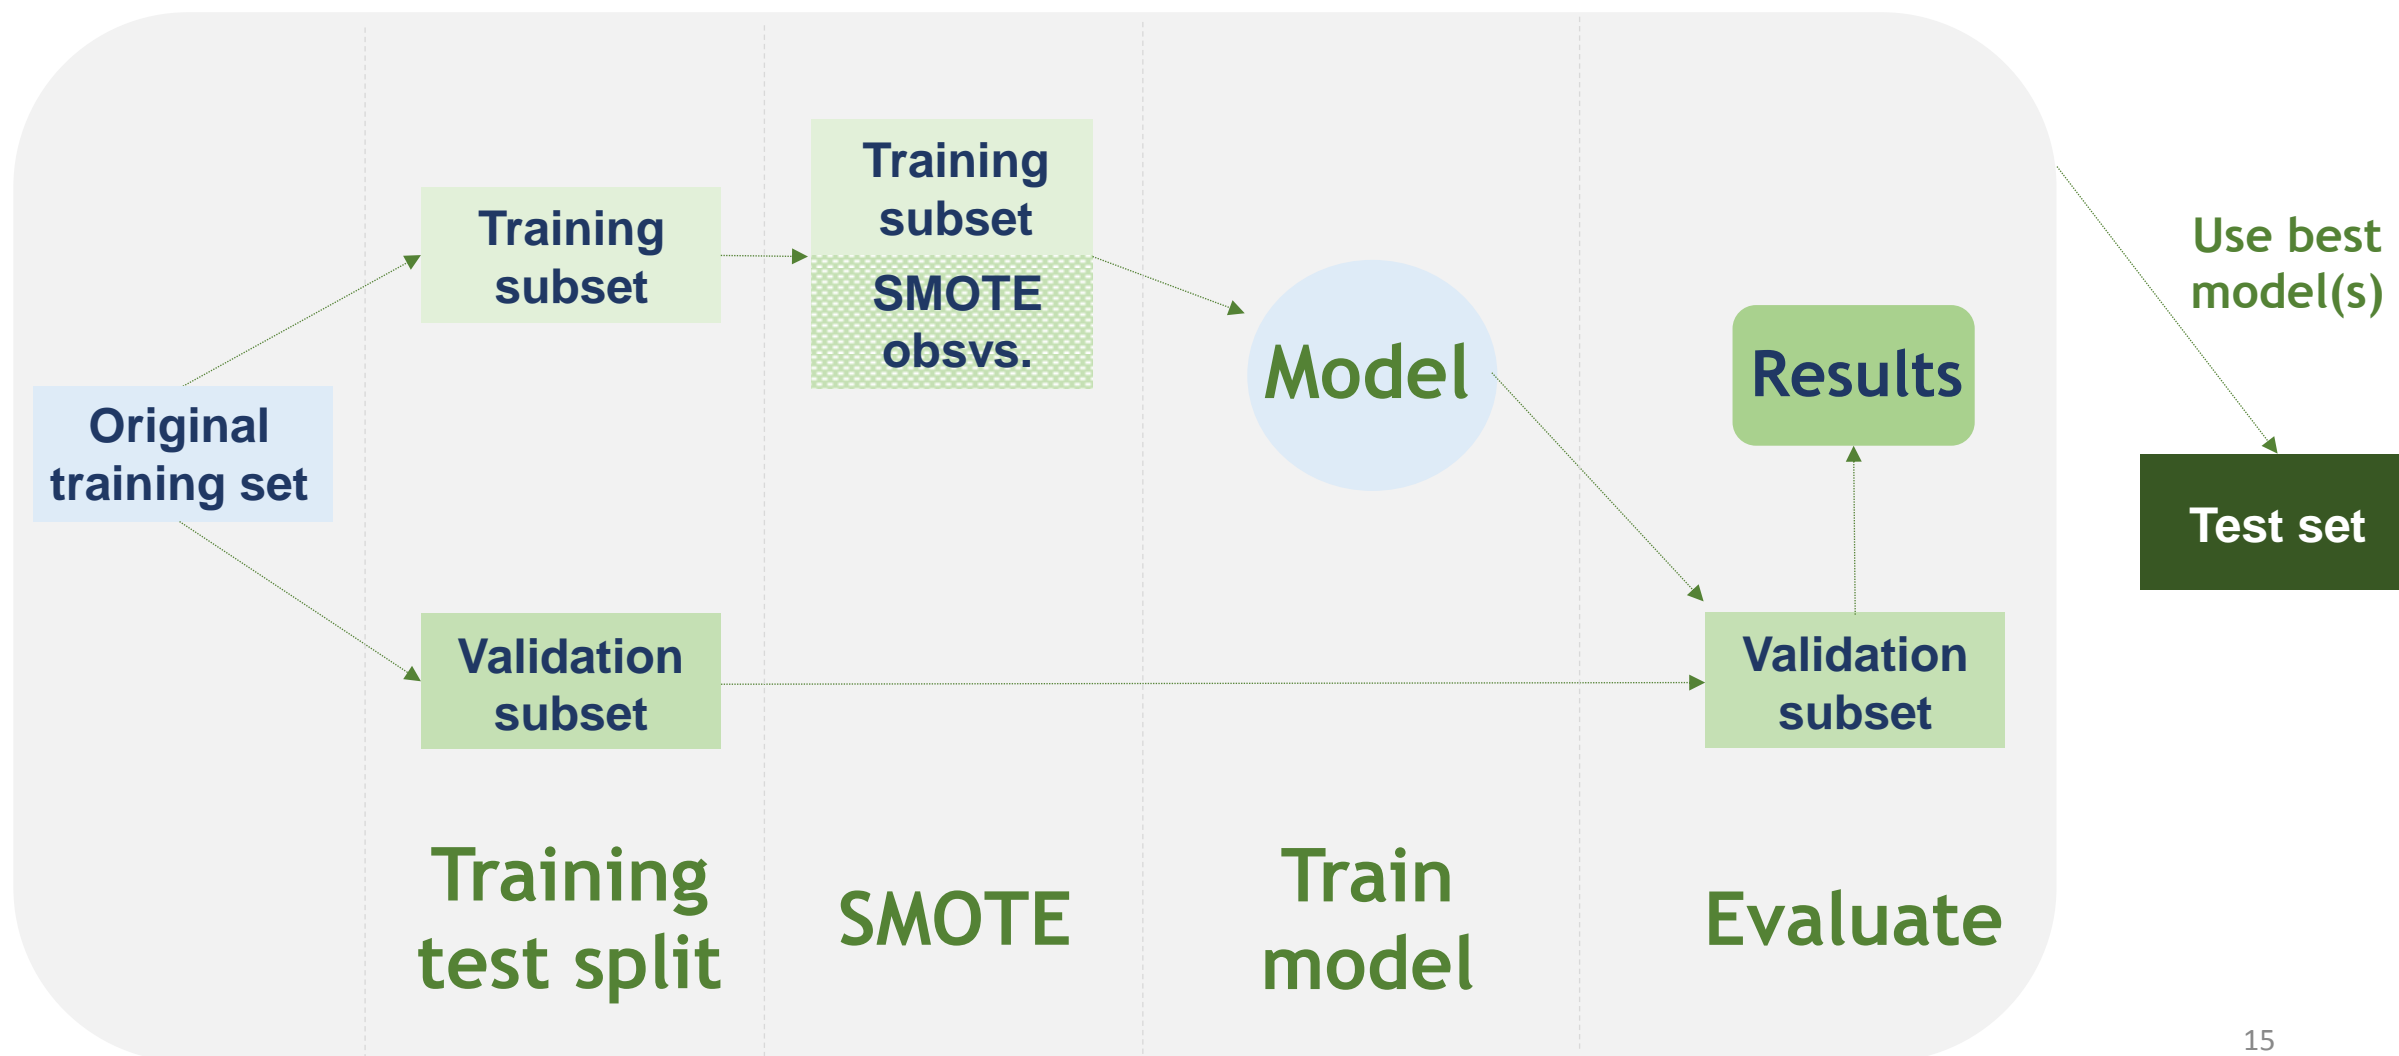


Majority class samples
Minority class samples
Synthetic samples


MODEL PIPELINE AND EVALUATION



Model pipeline



MODEL PIPELINE AND EVALUATION




Model performance evaluation metrics: a purely statistical approach

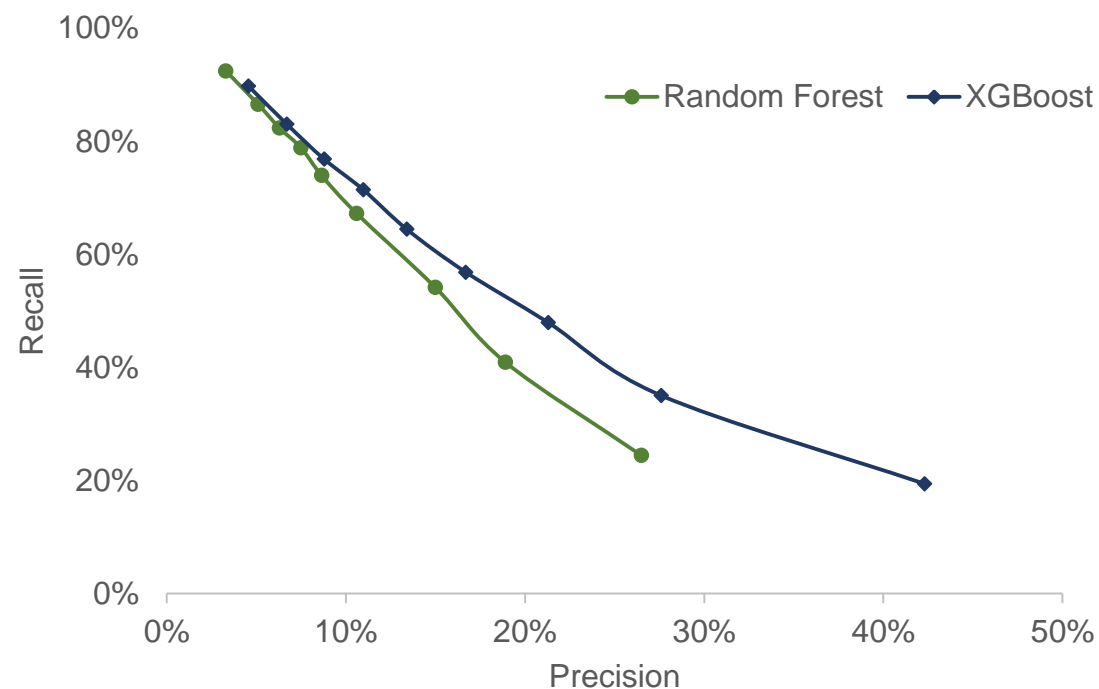

- In our opinion, accuracy is a poor measure of performance in a classification task with such a severe imbalance*
- We focus mostly on
 - Recall: percentage of back orders correctly predicted
 - Area under the ROC curve (AUC): probability that a positive record is given a higher likelihood of being a back order than a negative record
- We also monitor precision to get a sense of how many incorrectly predicted back orders our model produces

*Predicting that all records will equal 0 yields an accuracy of 99%


MODEL PIPELINE AND EVALUATION



Model performance evaluation metrics: a purely statistical approach



MODEL PIPELINE AND EVALUATION




Model performance evaluation metrics: a business cost approach

	Pred 0	Pred 1
True 0	344,318	49,697
True 1	295	2,432

Recall = 89%
Precision = 5%

- Is a model any good if it correctly predicts 89% of all true back orders, but only 5% of the model's predicted back orders become true?
- If the company faces a cost when incorrectly predicting a back order (e.g., because it is stocking up on inventory that it doesn't need), then a model with 5% precision might be highly undesirable, depending on the cost
- From the business case provided, we cannot know such costs, but we can try to rank models accordingly

MODEL PIPELINE AND EVALUATION




Model performance evaluation metrics: a business cost approach

	Pred 0	Pred 1
True 0	344,318	49,697
True 1	295	2,432

Suppose:

- The cost of a False Positive (FP) is x
 - And the cost of a False Negative (FN) is y
 - The model would produce a cost of $295y + 49,697x$
-
- To simplify, we assume the baseline case is that the firm does not try to predict back orders at the moment. Therefore, it faces a cost equal to $(295 + 2,432)y$
 - The firm should adopt the model if $295y + 49,697x < (295 + 2,432)y$

MODEL PIPELINE AND EVALUATION




Model performance evaluation metrics: a business cost approach

- More generally, the firm should adopt the model when $(FN + TP) \cdot y > FN \cdot y + FP \cdot x$, i.e., when the baseline cost is greater than the model cost
- This is the same result as


$$\frac{FP}{TP} < \frac{y}{x}$$

Which means that the firm should always adopt a model when the ratio FP/TP is smaller than the ratio of back order costs (y) to excess inventory costs (x). Therefore, in addition to recall and AUC, we will favor models with low FP / TP (let us call it FPTP ratio)


MODEL PIPELINE AND EVALUATION



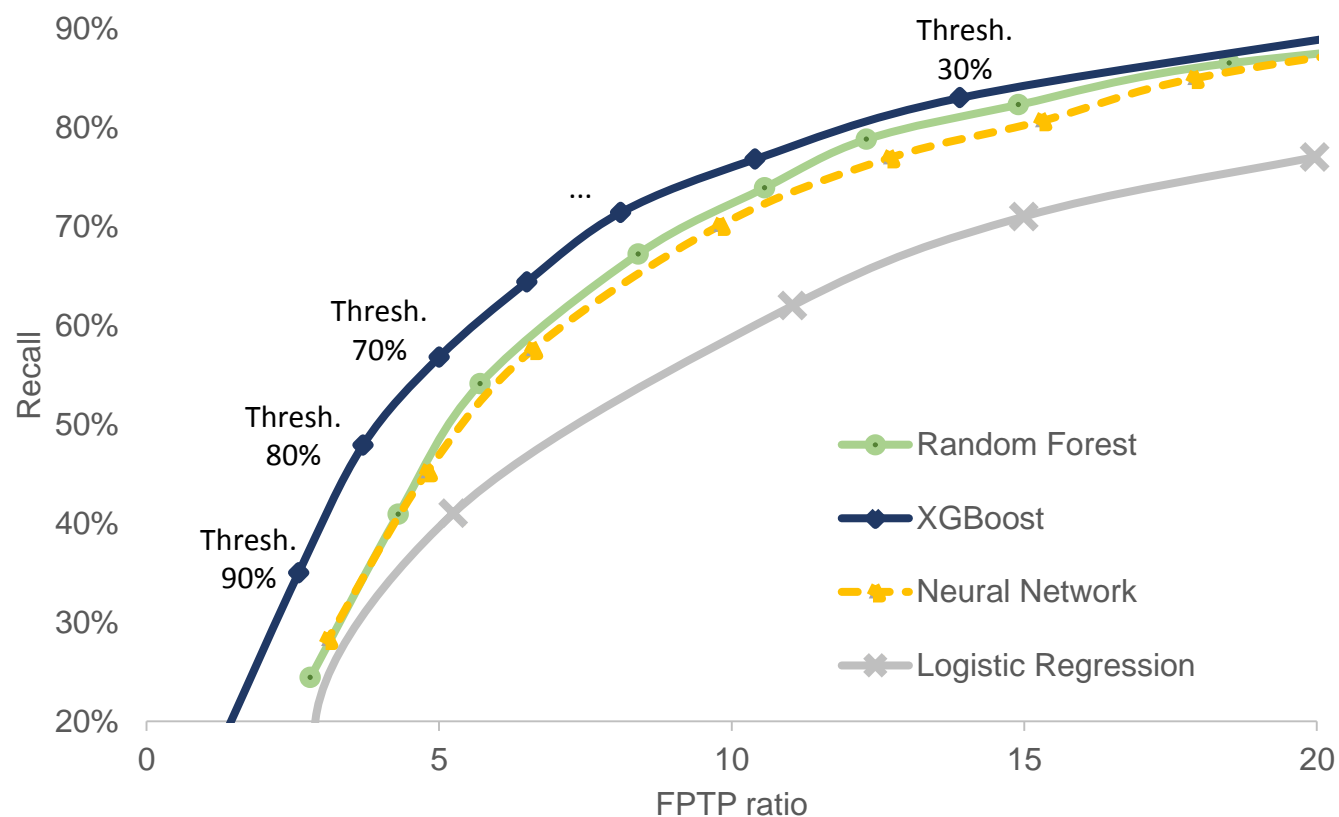
Model performance evaluation metrics: a business cost approach



MODEL PIPELINE AND EVALUATION



Model performance evaluation



Model	Recall	AUC	Precision	FPTP
Random forest with 1,000 trees	74%	84%	9%	10,6
Balanced RF with 1,000 trees	89%	86%	4%	20,4
KNN with 3 neighbors	43%	76%	16%	5,1
Logistic Regression	72%	82%	6%	14,9
Neural Network	75%	94%	8%	12,1
XGBoost with 1,000 trees	64%	81%	13%	6,5

MODEL RESULTS AND PROOF OF CONCEPT

Model performance in predicting the Test set

- How would our model stack up in predicting the following week?
- Recall = 67%, FPTP ratio = 9.2, AUC = 80%
- Perhaps more importantly, **how much would the company save?**

	Pred 0	Pred 1
True 0	208,683	16,064
True 1	850	1,754



- For the sake of the example, suppose the expected cost of a backorder is 10€ and the cost of a False Positive is 1€ (e.g. stocking up on too much inventory)
- Cost without predictions = $(850 + 1,754) \times 10\text{€} = \mathbf{26,040\text{€}}$
- Costs when using the model = $850 \times 10\text{€} + 16,064 \times 1\text{€} = \mathbf{16,914\text{€}}$
- The company would have saved 9,126€ in 1 week, which in annual terms would imply a cost saving above 470,000€

CONCLUSIONS



- In a prediction problem with imbalanced data, accuracy is a poor performance measure, and a cost-benefit analysis of the outcomes of our model can be more insightful
- Taking a more statistical approach, we would be inclined to say that our Neural Network model is our best model (high AUC)
- However, taking into account potential business costs, the NN model might be too “imprecise”, and we’d propose using the XGBoost model, which could be more profitable for the firm

THANK YOU

ANNEX: ASSESSING FEATURE IMPORTANCE

- In order to use only relevant features, we calculate feature importance in a random forest model through the so called “Gini importance”, which is measured as is defined as the total decrease in node impurity (weighted by the probability of reaching that node) averaged over all trees of the ensemble.
- The Gini importance essentially evaluates the decrease in model accuracy when one randomly permutes the values for a feature. The higher the decrease, the greater the importance in that feature. A higher level of the measure indicates the feature is more relevant.

	importance
national_inv	3.54
forecast_3_month	1.43
forecast_6_month	1.08
forecast_9_month	0.58
in_transit_qty	0.50
sales_1_month	0.50
sales_3_month	0.41
stop_auto_buy	0.30
local_bo_qty	0.28
sales_6_month	0.24
sales_9_month	0.21
lead_time	0.19
perf_6_month_avg	0.19
perf_12_month_avg	0.18
min_bank	0.15
pieces_past_due	0.09
ppap_risk	0.08
deck_risk	0.07
potential_issue	0.00
oe_constraint	0.00
rev_stop	0.00

ANNEX: CHOOSING AN APPROPRIATE SMOTE RATIO

Model	Recall	AUC	Precision	FPTP ratio
Random forest with 100 trees, SMOTE = 0%	0%	50%	100%	0,0
Random forest with 100 trees, SMOTE = 3%	24%	63%	30%	2,3
Random forest with 100 trees, SMOTE = 5%	36%	67%	23%	3,3
Random forest with 100 trees, SMOTE = 10%	48%	73%	18%	4,6
Random forest with 100 trees, SMOTE = 20%	67%	82%	11%	8,4
Random forest with 100 trees, SMOTE = 30%	74%	84%	9%	10,5
Random forest with 100 trees, SMOTE = 40%	77%	85%	8%	11,8
Random forest with 100 trees, SMOTE = 100%	85%	87%	6%	17,0

- The table above shows the KPIs in a Random Forest model for different levels of the SMOTE ratio
- Given the trade-off that we believe exists between Recall and FPTP, we concluded that the most balanced performance is achieved in the 30-40% proportion of minority class samples (generated using SMOTE)
- We decided to use 30% since it achieves similar performance and is computationally “less expensive” than 40%

ANNEX: PARAMETERIZING XGBOOST

Screenshot 1

```
params = {  
    'gamma': [1, 1.5, 2],  
    'subsample': [0.6, 0.8, 1.0],  
    'colsample_bytree': [0.6, 0.8, 1.0],  
    'learning_rate': [0.01, 0.02, 0.03, 0.04],  
}
```

```
xgb = XGBClassifier(n_estimators=600, max_depth=10, min_child_weight=5, objective='binary:logistic',  
                  silent=True, nthread=1)
```

```
folds = 5  
param_comb = 4
```

```
skf = StratifiedKFold(n_splits=folds, shuffle = True, random_state = 1001)
```

```
random_search = RandomizedSearchCV(xgb, param_distributions=params, n_iter=param_comb,  
                                   scoring='roc_auc', n_jobs=2, cv=skf.split(X_train_res, y_train_res),  
                                   verbose=3, random_state=1001 )
```

We chose XGBoost's hyperparameters through cross-validated randomized search, optimizing for AUC and with the initial set of parameters shown in screenshot 1

Screenshot 2

```
random_search.best_estimator_
```

```
XGBClassifier(base_score=0.5, booster='gbtree', colsample_bylevel=1,  
             colsample_bytree=0.6, gamma=1, learning_rate=0.03, max_delta_step=0,  
             max_depth=10, min_child_weight=5, missing=None, n_estimators=600,  
             n_jobs=1, nthread=1, objective='binary:logistic', random_state=0,  
             reg_alpha=0, reg_lambda=1, scale_pos_weight=1, seed=None,  
             silent=True, subsample=0.8)
```

The best parameters are shown in screenshot 2